

DnB Loan Default Prediction Assignment

August 10, 2022

1 Importing the Data Wrangling and Visualization libraries

```
[57]: import numpy as np
import pandas as pd

import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
sns.set_theme(style = "darkgrid")
```

2 Importing the dataset

```
[3]: df = pd.read_csv('/content/Borrower_Payment_History.csv')
df.head()
```

```
[3]:
```

	subject_id	account_id	DATE_REPORTED	currency_code	\
0	110113010000498773	2201160002396407	02-JAN-18	TZS	
1	110115010001596750	2201170002990772	08-JAN-18	TZS	
2	110114010000832340	2201140001177207	02-JAN-18	TZS	
3	110114010000740514	2201160002371670	02-JAN-18	TZS	
4	110114010000811717	2201140001161294	08-JAN-18	TZS	

	REPAYMENT_FREQUENCY	NUMBER_OF_INSTALLMENTS	SANCTION_AMT	\
0	MonthlyInstalments30Days	6.0	400000.0	
1	MonthlyInstalments30Days	61.0	3500000.0	
2	MonthlyInstalments30Days	57.0	4200000.0	
3	MonthlyInstalments30Days	24.0	60000000.0	
4	MonthlyInstalments30Days	69.0	3400000.0	

	TOT_OUTSTD_BAL	OVER_DUE_AMT	AMOUNT_OF_INSTALLMENT	...	LAST_AMOUNT_PAID	\
0	0.00	205609.4	0.00	...	NaN	
1	3272849.49	0.0	100243.11	...	NaN	
2	870462.93	0.0	106652.40	...	NaN	
3	15740676.93	0.0	2987342.81	...	NaN	
4	828981.51	0.0	99164.00	...	NaN	

	DATE_ACC_OPEN	DATE_ACC_CLOSE	INSTALMENT_LOAN_TYPE	loan_status	\
0	29-JUN-16	29-DEC-16	OtherInstalmentOperation	Existing	
1	29-JUN-17	05-JUL-22	ConsumerLoan	Existing	
2	27-DEC-13	NaN	ConsumerLoan	Existing	
3	11-MAY-16	NaN	BusinessLoan	Existing	
4	20-FEB-13	16-NOV-18	ConsumerLoan	Existing	

	ASSET_CLASSIFICATION	REASON_FOR_CLOSURE	NO_OF_DAYS_PAST_DUE	LOAN_TENURE	\
0	NaN	NaN	399.0	NaN	
1	NaN	NaN	0.0	NaN	
2	NaN	NaN	0.0	NaN	
3	NaN	NaN	18.0	NaN	
4	NaN	NaN	0.0	NaN	

	EXPECTED_END_DATE_OF_LOAN
0	29-DEC-16
1	05-JUL-22
2	27-SEP-18
3	11-MAY-18
4	16-NOV-18

[5 rows x 21 columns]

3 Explore the basics of the Dataset

```
[4]: # Check the size of the dataframe

rows, cols = df.shape
print('Rows: ', rows, '| Columns: ', cols)
```

Rows: 6000 | Columns: 21

```
[5]: # Understand the fields in the dataframe

df.columns
```

```
[5]: Index(['subject_id', 'account_id', 'DATE_REPORTED', 'currency_code',
            'REPAYMENT_FREQUENCY', 'NUMBER_OF_INSTALLMENTS', 'SANCTION_AMT',
            'TOT_OUTSTD_BAL', 'OVER_DUE_AMT', 'AMOUNT_OF_INSTALLMENT',
            'DATE_LATEST_PAY_RECEIVED', 'LAST_AMOUNT_PAID', 'DATE_ACC_OPEN',
            'DATE_ACC_CLOSE', 'INSTALMENT_LOAN_TYPE', 'loan_status',
            'ASSET_CLASSIFICATION', 'REASON_FOR_CLOSURE', 'NO_OF_DAYS_PAST_DUE',
            'LOAN_TENURE', 'EXPECTED_END_DATE_OF_LOAN'],
           dtype='object')
```

```
[6]: # Check the datatypes
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6000 entries, 0 to 5999
Data columns (total 21 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   subject_id                           6000 non-null   int64
1   account_id                           6000 non-null   int64
2   DATE_REPORTED                        6000 non-null   object
3   currency_code                        6000 non-null   object
4   REPAYMENT_FREQUENCY                  5982 non-null   object
5   NUMBER_OF_INSTALLMENTS               5921 non-null   float64
6   SANCTION_AMT                         6000 non-null   float64
7   TOT_OUTSTD_BAL                       5981 non-null   float64
8   OVER_DUE_AMT                        5993 non-null   float64
9   AMOUNT_OF_INSTALLMENT                5972 non-null   float64
10  DATE_LATEST_PAY_RECEIVED              4647 non-null   object
11  LAST_AMOUNT_PAID                     0 non-null      float64
12  DATE_ACC_OPEN                        6000 non-null   object
13  DATE_ACC_CLOSE                       2934 non-null   object
14  INSTALMENT_LOAN_TYPE                 5920 non-null   object
15  loan_status                          6000 non-null   object
16  ASSET_CLASSIFICATION                  0 non-null      float64
17  REASON_FOR_CLOSURE                    0 non-null      float64
18  NO_OF_DAYS_PAST_DUE                  5929 non-null   float64
19  LOAN_TENURE                           0 non-null      float64
20  EXPECTED_END_DATE_OF_LOAN            6000 non-null   object
dtypes: float64(10), int64(2), object(9)
memory usage: 984.5+ KB
```

4 Data Cleaning

```
[7]: # Drop the fields which are completely NULL
```

```
df1 = df.drop(columns=['LAST_AMOUNT_PAID', 'ASSET_CLASSIFICATION',
↳ 'REASON_FOR_CLOSURE', 'LOAN_TENURE'])
df1.shape
```

```
[7]: (6000, 17)
```

```
[8]: # Drop duplicated records from the dataframe
```

```
df1 = df1.drop_duplicates()
df1.shape
```

[8]: (5982, 17)

[9]: *# Rename the fields uniformly.*

```
df1.rename(columns={'subject_id': 'SUBJECT_ID', 'account_id': 'ACCOUNT_ID',
                    'currency_code': 'CURR_CODE', 'INSTALMENT_LOAN_TYPE': 'LOAN_TYPE',
                    'AMOUNT_OF_INSTALLMENT': 'INSTALLMENT_AMT',
                    'DATE_LATEST_PAY_RECEIVED': 'DATE_LATEST_PAYMENT',
                    'loan_status': 'LOAN_STATUS', 'NO_OF_DAYS_PAST_DUE': 'DPD',
                    'EXPECTED_END_DATE_OF_LOAN': 'DATE_EXP_LOAN_CLOSURE'},
           inplace=True)
```

[10]: *# Correct the date fields to the right formats*

```
df1['DATE_REPORTED'] = pd.to_datetime(df1['DATE_REPORTED'])
df1['DATE_ACC_OPEN'] = pd.to_datetime(df1['DATE_ACC_OPEN'])
df1['DATE_ACC_CLOSE'] = pd.to_datetime(df1['DATE_ACC_CLOSE'])
df1['DATE_EXP_LOAN_CLOSURE'] = pd.to_datetime(df1['DATE_EXP_LOAN_CLOSURE'])
df1['DATE_LATEST_PAYMENT'] = pd.to_datetime(df1['DATE_LATEST_PAYMENT'])

df1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 5982 entries, 0 to 5999
Data columns (total 17 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   SUBJECT_ID                            5982 non-null   int64
1   ACCOUNT_ID                            5982 non-null   int64
2   DATE_REPORTED                         5982 non-null   datetime64[ns]
3   CURR_CODE                             5982 non-null   object
4   REPAYMENT_FREQUENCY                   5964 non-null   object
5   NUMBER_OF_INSTALLMENTS                 5903 non-null   float64
6   SANCTION_AMT                          5982 non-null   float64
7   TOT_OUTSTD_BAL                        5963 non-null   float64
8   OVER_DUE_AMT                          5975 non-null   float64
9   INSTALLMENT_AMT                       5954 non-null   float64
10  DATE_LATEST_PAYMENT                    4631 non-null   datetime64[ns]
11  DATE_ACC_OPEN                          5982 non-null   datetime64[ns]
12  DATE_ACC_CLOSE                         2927 non-null   datetime64[ns]
13  LOAN_TYPE                              5902 non-null   object
14  LOAN_STATUS                            5982 non-null   object
15  DPD                                    5911 non-null   float64
```

```

16 DATE_EXP_LOAN_CLOSURE    5982 non-null    datetime64[ns]
dtypes: datetime64[ns](5), float64(6), int64(2), object(4)
memory usage: 841.2+ KB

```

```
[11]: # Note that the Account_ID is repeatedly reported on multiple date.
```

```
df1[df1['ACCOUNT_ID']== 2201170002918195]
```

```
[11]:
```

	SUBJECT_ID	ACCOUNT_ID	DATE_REPORTED	CURR_CODE	\
137	110117010002384406	2201170002918195	2018-03-08	TZS	
4369	110117010002384406	2201170002918195	2018-06-10	TZS	

	REPAYMENT_FREQUENCY	NUMBER_OF_INSTALLMENTS	SANCTION_AMT	\
137	MonthlyInstalments30Days	61.0	10200000.0	
4369	MonthlyInstalments30Days	61.0	10200000.0	

	TOT_OUTSTD_BAL	OVER_DUE_AMT	INSTALLMENT_AMT	DATE_LATEST_PAYMENT	\
137	8192765.61	0.0	270237.61	2018-02-27	
4369	7787140.81	0.0	270237.61	2018-05-25	

	DATE_ACC_OPEN	DATE_ACC_CLOSE	LOAN_TYPE	LOAN_STATUS	DPD	\
137	2016-08-19	2021-09-01	ConsumerLoan	Existing	0.0	
4369	2016-08-19	2021-09-01	ConsumerLoan	Existing	0.0	

	DATE_EXP_LOAN_CLOSURE
137	2021-09-01
4369	2021-09-01

```
[12]: # Remove multiple entries for same Account_ID reported on different dates to
↳ make it unique.
```

```
df2 = df1.merge(df1.groupby('ACCOUNT_ID')['DATE_REPORTED'].max().reset_index(),
                on=['ACCOUNT_ID', 'DATE_REPORTED'], how='inner')
df2.shape
```

```
[12]: (5969, 17)
```

```
[13]: # Drop the records with SANCTION_AMT = 0, which doesnt add value to dataset.
```

```
df2.drop(df2[df2['SANCTION_AMT'] == 0].index, inplace = True)
df2.shape
```

```
[13]: (5960, 17)
```

5 Handle Missing Values

```
[14]: # Check the NULL values in the dataframe
```

```
df2.isnull().sum()
```

```
[14]: SUBJECT_ID          0
      ACCOUNT_ID         0
      DATE_REPORTED      0
      CURR_CODE          0
      REPAYMENT_FREQUENCY 12
      NUMBER_OF_INSTALLMENTS 70
      SANCTION_AMT       0
      TOT_OUTSTD_BAL     13
      OVER_DUE_AMT       7
      INSTALLMENT_AMT    22
      DATE_LATEST_PAYMENT 1345
      DATE_ACC_OPEN      0
      DATE_ACC_CLOSE     3046
      LOAN_TYPE          80
      LOAN_STATUS        0
      DPD               71
      DATE_EXP_LOAN_CLOSURE 0
      dtype: int64
```

```
[15]: # Find the values in REPAYMENT_FREQUENCY.
```

```
df2['REPAYMENT_FREQUENCY'].value_counts()
```

```
[15]: MonthlyInstalments30Days      5684
      AtTheFinalDayOfThePeriodOfContract    197
      IrregularInstalments              40
      FortnightlyInstalments15Days        25
      AnnualInstalments360Days            1
      SixMonthInstalments180Days          1
      Name: REPAYMENT_FREQUENCY, dtype: int64
```

```
[16]: # Find the values in LOAN_TYPE.
```

```
df2['LOAN_TYPE'].value_counts()
```

```
[16]: ConsumerLoan      3795
      BusinessLoan   1678
      OtherInstalmentOperation    310
      CreditCard      58
      MortgageLoan    16
      Overdraft       12
```

```
LeasingFinancial          11
Name: LOAN_TYPE, dtype: int64
```

```
[17]: # Find the values in LOAN_STATUS.
```

```
df2['LOAN_STATUS'].value_counts()
```

```
[17]: Existing          4914
TerminatedAccordingTheContract    760
TerminatedInAdvanceCorrectly      279
TerminatedInAdvanceIncorrectly      7
Name: LOAN_STATUS, dtype: int64
```

```
[19]: # Fill the missing the 'REPAYMENT_FREQUENCY' field values
```

```
# For LOAN_TYPE == 'Overdraft', the payment cannot have installments but 1
↳ single payment.
```

```
df2['REPAYMENT_FREQUENCY'] = np.where(df2['LOAN_TYPE'] == "Overdraft",
                                      'AtTheFinalDayOfThePeriodOfContract',
                                      df2['REPAYMENT_FREQUENCY'])
```

```
# # For the rest, lets fill it with the mode of the field values.
```

```
# df2['REPAYMENT_FREQUENCY'].fillna(df2['REPAYMENT_FREQUENCY'].mode()[0],
↳ inplace=True)
```

```
[22]: # Create a function which calculates the number of installments based on the
↳ DATE_ACC_OPEN and DATE_EXP_LOAN_CLOSURE
```

```
def calc_instal_count(freq, end_date, open_date):
    if freq == 'FortnightlyInstalments15Days':
        tenure = (end_date - open_date) / np.timedelta64(1, 'D') / 15
        return tenure
    elif freq == 'MonthlyInstalments30Days':
        tenure = (end_date - open_date) / np.timedelta64(1, 'D') / 30
        return tenure
    elif freq == 'AnnualInstalments360Days':
        tenure = (end_date - open_date) / np.timedelta64(1, 'D') / 360
        return tenure
    elif freq == 'SixMonthInstalments180Days':
        tenure = (end_date - open_date) / np.timedelta64(1, 'D') / 180
        return tenure
    elif freq == 'AtTheFinalDayOfThePeriodOfContract':
        tenure = 1
        return tenure
    else:
        return 1
```

```
[23]: # Replace the Number of Installments field with calculated NON NULL values

df2['NUMBER_OF_INSTALLMENTS'] = df2.apply(lambda x:
    ↪ calc_instal_count(x['REPAYMENT_FREQUENCY'],
    ↪
    ↪ x['DATE_EXP_LOAN_CLOSURE'],
    ↪
    ↪ x['DATE_ACC_OPEN']), axis=1)
df2['NUMBER_OF_INSTALLMENTS'] = df2['NUMBER_OF_INSTALLMENTS'].apply(np.floor)

[25]: # Fill in the INSTALLMENT_AMT field by dividing SANCTION_AMT with
    ↪ NUMBER_OF_INSTALLMENTS

df2['INSTALLMENT_AMT'].fillna(df2['SANCTION_AMT'] /
    ↪ df2['NUMBER_OF_INSTALLMENTS'], inplace=True)

[29]: # Fill in missing the OVER_DUE_AMT

# For the loans which are Terminated properly, there should not be any Over-Due
    ↪ Amount
df2.loc[df2['LOAN_STATUS']=='TerminatedAccordingTheContract', 'OVER_DUE_AMT'] = 0

# The loans with DPD ==0, there should not be Over_Due_Amount
df2.loc[(df2['OVER_DUE_AMT'].isnull()) & (df2['DPD']==0), 'OVER_DUE_AMT'] = 0

# For the rest of the loans with status "Existing", since they all have
    ↪ Repayment Frequency of 30 days &
# since the DPDs are less than 30, their OVER_DUE should be equal to
    ↪ INSTALLMENT_AMT
df2['OVER_DUE_AMT'].fillna(df2['INSTALLMENT_AMT'], inplace=True)

[36]: # Fill in the TOT-OUTSTD_BAL field with 2 logics.

# For DPD >0 customers, their Oust_Bal must be greater than Over_Due.
# However, we shall consider the same amount here as we dont have the
    ↪ Total_Paid on date.
df2.loc[(df2['TOT_OUTSTD_BAL'].isnull()) & (df2['DPD']>0), 'TOT_OUTSTD_BAL'] =
    ↪ df2['OVER_DUE_AMT']

# For DPD ==0 customers, their outstanding is at the least their next
    ↪ INSTALLMENT_AMT
df2.loc[(df2['TOT_OUTSTD_BAL'].isnull()) & (df2['DPD']==0), 'TOT_OUTSTD_BAL'] =
    ↪ df2['INSTALLMENT_AMT']

[39]: # Fill in the missing LOAN_TYPE with the mode value.
```



```
# Because, find the the correct Loan_type using the Number of Installments
↳ seems difficult
# as even the loans from Credit Cards are allowed to pay in longer installments

df2['LOAN_TYPE'].fillna(df2['LOAN_TYPE'].mode()[0], inplace=True)
```

```
[40]: # Fill in the missing DPD days.

# For those loans which have valid LATEST_PAYMENT_DATE, if the OVER_DUE exists,
↳ find the days_diff from reporting date.
df2.loc[(df2['OVER_DUE_AMT']>0) &
        (df2['DATE_LATEST_PAYMENT'].notnull()) &
        (df2['DPD'].isnull()), 'DPD'] = (df2['DATE_REPORTED'] -
↳ df2['DATE_LATEST_PAYMENT']) / np.timedelta64(1, 'D')
```

```
[41]: # Loans which are to be paid at the end of the contract, DPD should be 0
df2.loc[(df2['DPD'].isnull()) &
        (df2['REPAYMENT_FREQUENCY']=='AtTheFinalDayOfThePeriodOfContract'),'DPD'] = 0

# Notice that for the rest records with NULL DPD, REPAYMENY_FREQ is every month.
# So, if the Overdue_Amount is greater than 3 times of the installment, it has
↳ DPD > 90
df2.loc[(df2['DPD'].isnull()) &
        ((df2['OVER_DUE_AMT']/df2['INSTALLMENT_AMT'])>=3), 'DPD'] = 91

# Else, it has lesser than 90 days DPD. Fill the rest NULL DPD with 0.
df2['DPD'].fillna(0, inplace=True)
```

```
[43]: # Thus cleaned all the NULL values on necessary fields.

df2.isnull().sum()
```

```
[43]: SUBJECT_ID          0
ACCOUNT_ID             0
DATE_REPORTED          0
CURR_CODE              0
REPAYMENT_FREQUENCY    0
NUMBER_OF_INSTALLMENTS 0
SANCTION_AMT           0
TOT_OUTSTD_BAL         0
OVER_DUE_AMT           0
INSTALLMENT_AMT        0
DATE_LATEST_PAYMENT    1345
DATE_ACC_OPEN          0
DATE_ACC_CLOSE         3046
LOAN_TYPE              0
```

```

LOAN_STATUS          0
DPD                  0
DATE_EXP_LOAN_CLOSURE 0
dtype: int64

```

```

[44]: # Adding a new feature 'LOAN_TENURE', which originally had all NULL values in
      ↪ the dataset.

df2['LOAN_TENURE'] = (df2['DATE_EXP_LOAN_CLOSURE'] - df2['DATE_ACC_OPEN']) / np.
      ↪ timedelta64(1, 'D')
df2.head()

```

```

[44]:
      SUBJECT_ID      ACCOUNT_ID DATE_REPORTED  CURR_CODE  \
0  110113010000498773  2201160002396407    2018-01-02    TZS
1  110115010001596750  2201170002990772    2018-01-08    TZS
2  110114010000832340  2201140001177207    2018-01-02    TZS
3  110114010000740514  2201160002371670    2018-01-02    TZS
4  110114010000811717  2201140001161294    2018-01-08    TZS

      REPAYMENT_FREQUENCY  NUMBER_OF_INSTALLMENTS  SANCTION_AMT  \
0  MonthlyInstalments30Days                6.0      400000.0
1  MonthlyInstalments30Days               61.0     3500000.0
2  MonthlyInstalments30Days               57.0     4200000.0
3  MonthlyInstalments30Days               24.0     60000000.0
4  MonthlyInstalments30Days               69.0      3400000.0

      TOT_OUTSTD_BAL  OVER_DUE_AMT  INSTALLMENT_AMT  DATE_LATEST_PAYMENT  \
0              0.00     205609.4              0.00      2016-11-29
1      3272849.49              0.0      100243.11      2017-12-29
2      870462.93              0.0      106652.40      2017-12-28
3     15740676.93              0.0     2987342.81      2017-12-18
4      828981.51              0.0       99164.00      2017-12-29

      DATE_ACC_OPEN  DATE_ACC_CLOSE      LOAN_TYPE  LOAN_STATUS  DPD  \
0    2016-06-29    2016-12-29  OtherInstalmentOperation    Existing  399.0
1    2017-06-29    2022-07-05      ConsumerLoan    Existing    0.0
2    2013-12-27              NaT      ConsumerLoan    Existing    0.0
3    2016-05-11              NaT      BusinessLoan    Existing   18.0
4    2013-02-20    2018-11-16      ConsumerLoan    Existing    0.0

      DATE_EXP_LOAN_CLOSURE  LOAN_TENURE
0          2016-12-29          183.0
1          2022-07-05         1832.0
2          2018-09-27         1735.0
3          2018-05-11          730.0
4          2018-11-16         2095.0

```

```
[45]: # Select only the relevant attributes for further steps

features = ['SUBJECT_ID', 'ACCOUNT_ID', 'CURR_CODE', 'REPAYMENT_FREQUENCY',
            ↪ 'NUMBER_OF_INSTALLMENTS', 'SANCTION_AMT',
            'TOT_OUTSTD_BAL', 'OVER_DUE_AMT', 'INSTALLMENT_AMT', 'LOAN_TYPE',
            ↪ 'LOAN_STATUS', 'LOAN_TENURE', 'DPD']
df3 = df2[features]
df3.head()
```

```
[45]:
```

	SUBJECT_ID	ACCOUNT_ID	CURR_CODE	REPAYMENT_FREQUENCY	\
0	110113010000498773	2201160002396407	TZS	MonthlyInstalments30Days	
1	110115010001596750	2201170002990772	TZS	MonthlyInstalments30Days	
2	110114010000832340	2201140001177207	TZS	MonthlyInstalments30Days	
3	110114010000740514	2201160002371670	TZS	MonthlyInstalments30Days	
4	110114010000811717	2201140001161294	TZS	MonthlyInstalments30Days	

	NUMBER_OF_INSTALLMENTS	SANCTION_AMT	TOT_OUTSTD_BAL	OVER_DUE_AMT	\
0	6.0	400000.0	0.00	205609.4	
1	61.0	3500000.0	3272849.49	0.0	
2	57.0	4200000.0	870462.93	0.0	
3	24.0	60000000.0	15740676.93	0.0	
4	69.0	3400000.0	828981.51	0.0	

	INSTALLMENT_AMT	LOAN_TYPE	LOAN_STATUS	LOAN_TENURE	DPD
0	0.00	OtherInstalmentOperation	Existing	183.0	399.0
1	100243.11	ConsumerLoan	Existing	1832.0	0.0
2	106652.40	ConsumerLoan	Existing	1735.0	0.0
3	2987342.81	BusinessLoan	Existing	730.0	18.0
4	99164.00	ConsumerLoan	Existing	2095.0	0.0

```
[46]: df3.shape
```

```
[46]: (5960, 13)
```

6 Data Analysis & Visualisation

```
[47]: # Account_ID becomes the unique identifier for the dataframe.
```

```
df3['ACCOUNT_ID'].nunique()
```

```
[47]: 5960
```

```
[48]: # Note that some Subject_IDs are repeated in the dataframe, meaning some
      ↪ Customers availed more than 1 loan.
```

```
df3['SUBJECT_ID'].nunique()
```

[48]: 5923

```
[56]: # Check the customers who availed multiple loans and their default

df_multiple_loans = df3.groupby(['SUBJECT_ID'])['ACCOUNT_ID'].count().
    ↪sort_values(ascending=False).reset_index()
df_multiple_loans = df_multiple_loans[df_multiple_loans['ACCOUNT_ID'] > 1]
df_multiple_loans = df3[df3['SUBJECT_ID'].isin(df_multiple_loans['SUBJECT_ID'])]
df_multiple_loans.sort_values(by=['DPD', 'SUBJECT_ID'], ascending=False,
    ↪inplace=True)
df_multiple_loans[['SUBJECT_ID', 'ACCOUNT_ID', 'LOAN_TYPE', 'LOAN_STATUS',
    ↪'DPD']]
```

/usr/local/lib/python3.7/dist-packages/pandas/util/_decorators.py:311:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
return func(*args, **kwargs)
```

```
[56]:
```

	SUBJECT_ID	ACCOUNT_ID	LOAN_TYPE	LOAN_STATUS	DPD
588	110113010000508823	2201150001612310	ConsumerLoan	Existing	286.0
4451	110118010002814358	2201180003748597	BusinessLoan	Existing	0.0
5011	110118010002814358	2201180003748608	BusinessLoan	Existing	0.0
3537	110118010002810117	2201180003993041	BusinessLoan	Existing	0.0
4486	110118010002810117	2201180003719814	BusinessLoan	Existing	0.0
...
124	110113010000300324	2201180003234962	ConsumerLoan	Existing	0.0
5689	110113010000300324	2201180004039449	ConsumerLoan	Existing	0.0
2905	110113010000272956	2201180003748591	BusinessLoan	Existing	0.0
3227	110113010000272956	2201180003748599	BusinessLoan	Existing	0.0
3570	110113010000272956	2201180003869980	BusinessLoan	Existing	0.0

[73 rows x 5 columns]

```
[58]: df_multiple_loans['SUBJECT_ID'].nunique()
```

[58]: 36

6.1 Understanding the Categorical variables

```
[59]: # Function to plot Value Counts for Categorical variables
```

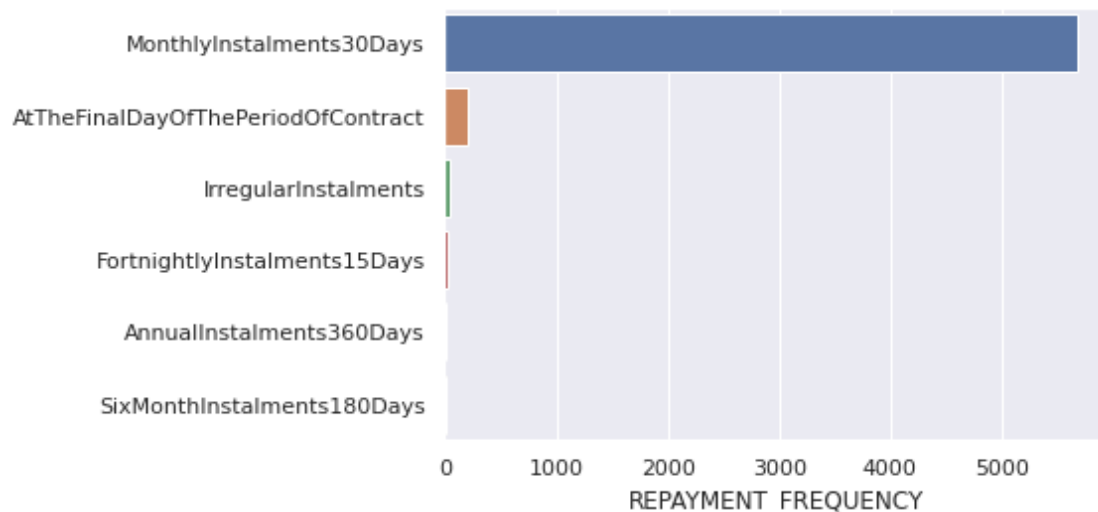
```
def cat_valcount_plot(df, fld):  
    plt.figure(figsize=(6, 4))  
    sns.barplot(df[fld].value_counts(), df[fld].value_counts().index)  
    plt.show()
```

```
[60]: # Check the distribution of multiple Repayment_Frequency
```

```
cat_valcount_plot(df3, 'REPAYMENT_FREQUENCY')
```

/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning

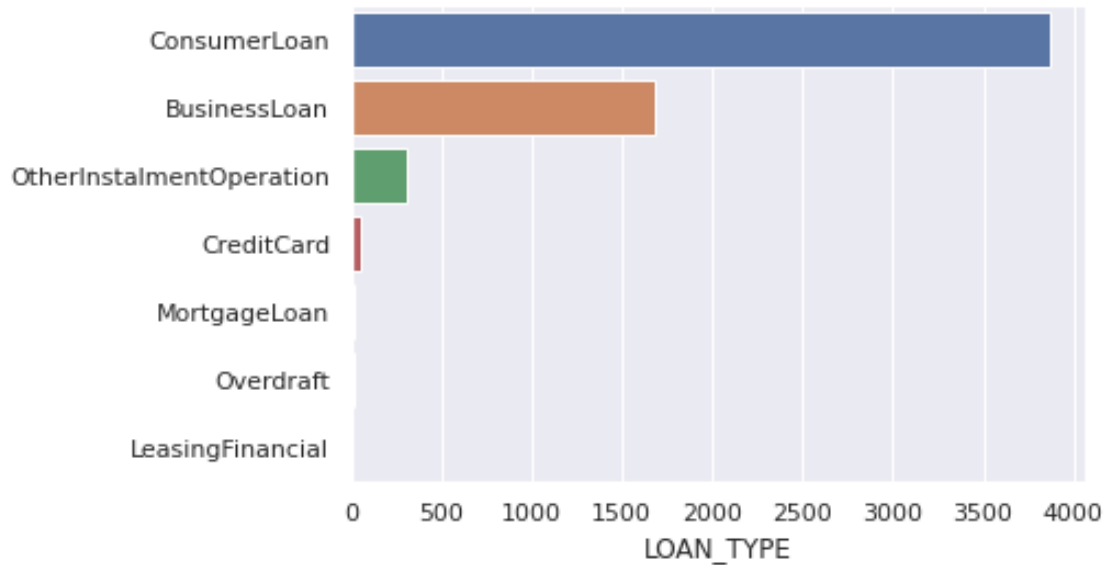


```
[61]: # Check the distribution of different Loan_Type
```

```
cat_valcount_plot(df3, 'LOAN_TYPE')
```

/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning

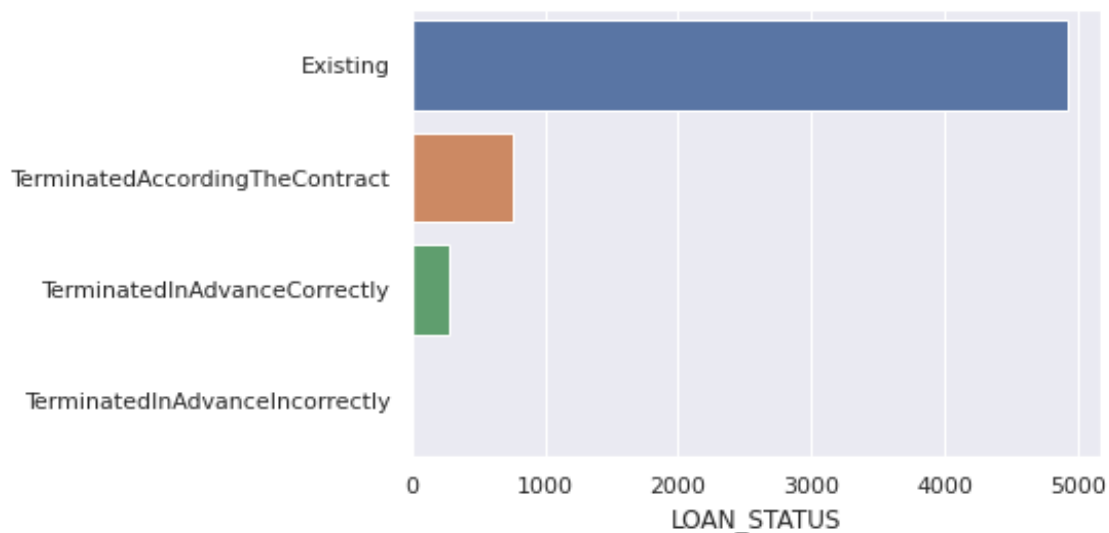


[62]: *# Check the distribution of different Loan_Status*

```
cat_valcount_plot(df3, 'LOAN_STATUS')
```

/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning




```
[63]: # Check the distribution of different Currencies
```

```
cat_valcount_plot(df3, 'CURR_CODE')
```

/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning



```
[64]: df3['CURR_CODE'].value_counts()
```

```
[64]: TZS    5955
      USD      5
      Name: CURR_CODE, dtype: int64
```

```
[65]: # Check the records which have currency in USD
```

```
df3[df3['CURR_CODE']=='USD'][['SANCTION_AMT', 'TOT_OUTSTD_BAL',
                              'OVER_DUE_AMT', 'INSTALLMENT_AMT']]
```

```
[65]:
```

	SANCTION_AMT	TOT_OUTSTD_BAL	OVER_DUE_AMT	INSTALLMENT_AMT
599	44500.0	37510.66	0.00	967.90
649	100000.0	0.00	0.00	775.00

902	60000.0	1484.13	1484.13	1484.13
4010	40000.0	15815.56	0.00	1309.55
4338	35000.0	33958.13	0.00	657.26

```
[66]: # Since there are only 5 Loans in USD, let's convert all USD Loans into TZS
```

```
# Exchange rate from USD TO TZS (today's rate is 2332)
usd_tzs_conv_rate = 2335

df3.loc[df3['CURR_CODE']=='USD', ['SANCTION_AMT', 'TOT_OUTSTD_BAL',
                                  'OVER_DUE_AMT', 'INSTALLMENT_AMT']] *= _
↳ usd_tzs_conv_rate
```

/usr/local/lib/python3.7/dist-packages/pandas/core/indexing.py:1884:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

self._setitem_single_column(loc, val, pi)

```
[67]: # Below USD values are now converted to TZS amounts
```

```
df3[df3['CURR_CODE']=='USD'][['SANCTION_AMT', 'TOT_OUTSTD_BAL',
                              'OVER_DUE_AMT', 'INSTALLMENT_AMT']]
```

```
[67]:
```

	SANCTION_AMT	TOT_OUTSTD_BAL	OVER_DUE_AMT	INSTALLMENT_AMT
599	103907500.0	87587391.10	0.00	2260046.50
649	233500000.0	0.00	0.00	1809625.00
902	140100000.0	3465443.55	3465443.55	3465443.55
4010	93400000.0	36929332.60	0.00	3057799.25
4338	81725000.0	79292233.55	0.00	1534702.10

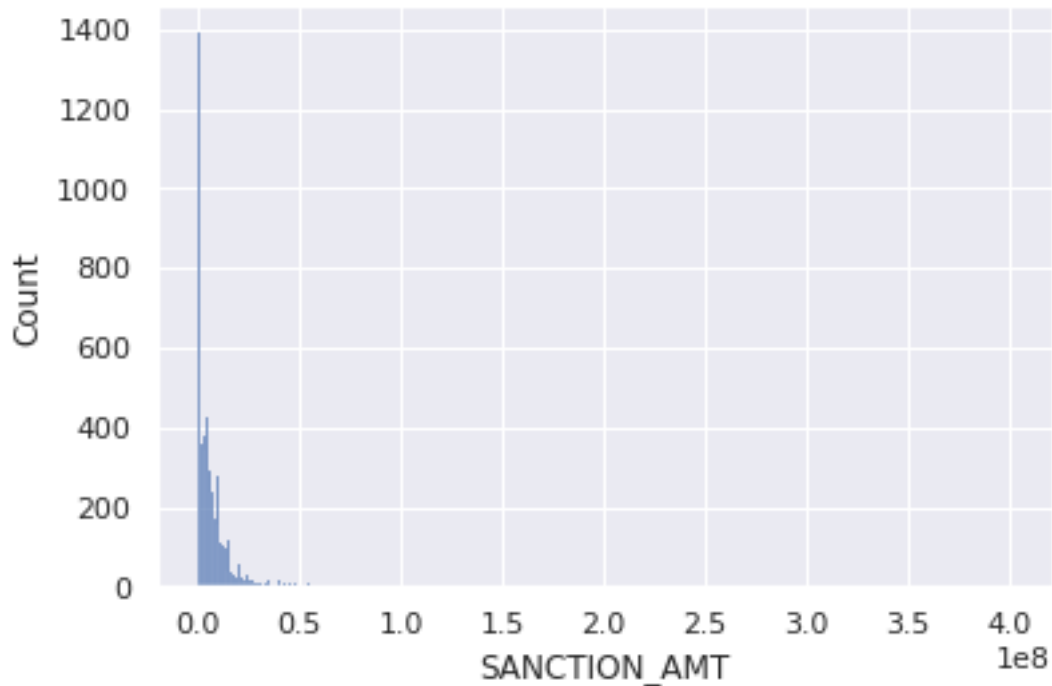
```
[68]: # Correcting the USD Currency code to TZS for the converted records
```

```
df3['CURR_CODE'].replace(['USD','TZS'], inplace=True)
```

7 Visualise & Normalise the Numerical Variables

```
[69]: sns.histplot(df3['SANCTION_AMT'])
```

```
[69]: <matplotlib.axes._subplots.AxesSubplot at 0x7fc29df51210>
```

```
[70]: # Applying log transformation to the attributes to normalise the data
```

```
df3['SANCTION_AMT_LOG'] = np.log(df3['SANCTION_AMT'])  
sns.histplot(df3['SANCTION_AMT_LOG'])
```

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:3:
```

```
SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame.
```

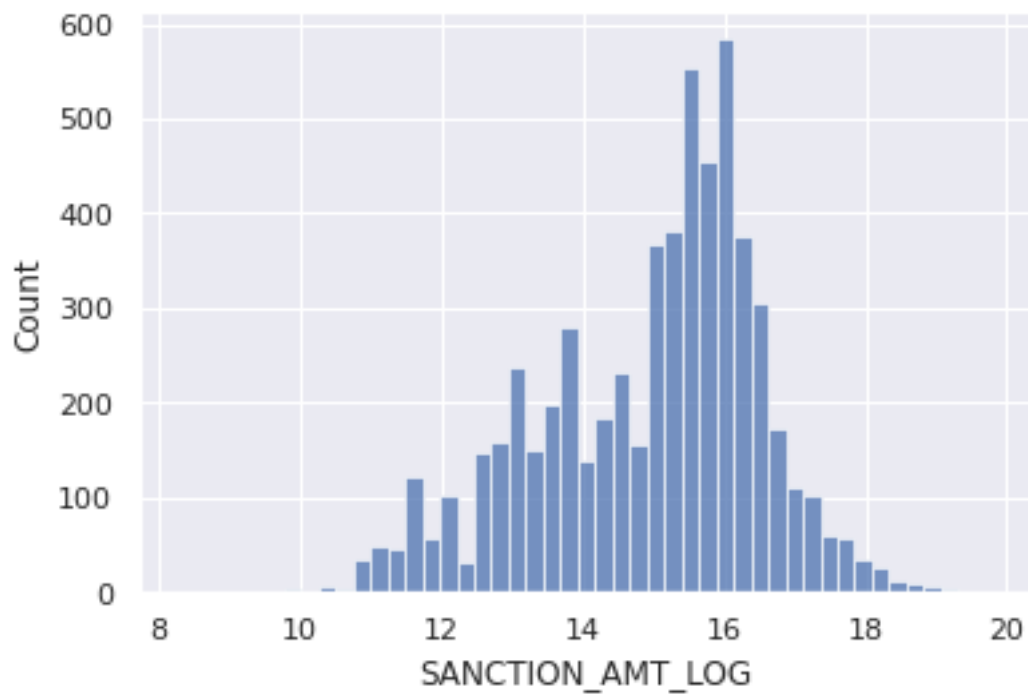
```
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-
```

```
docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
```

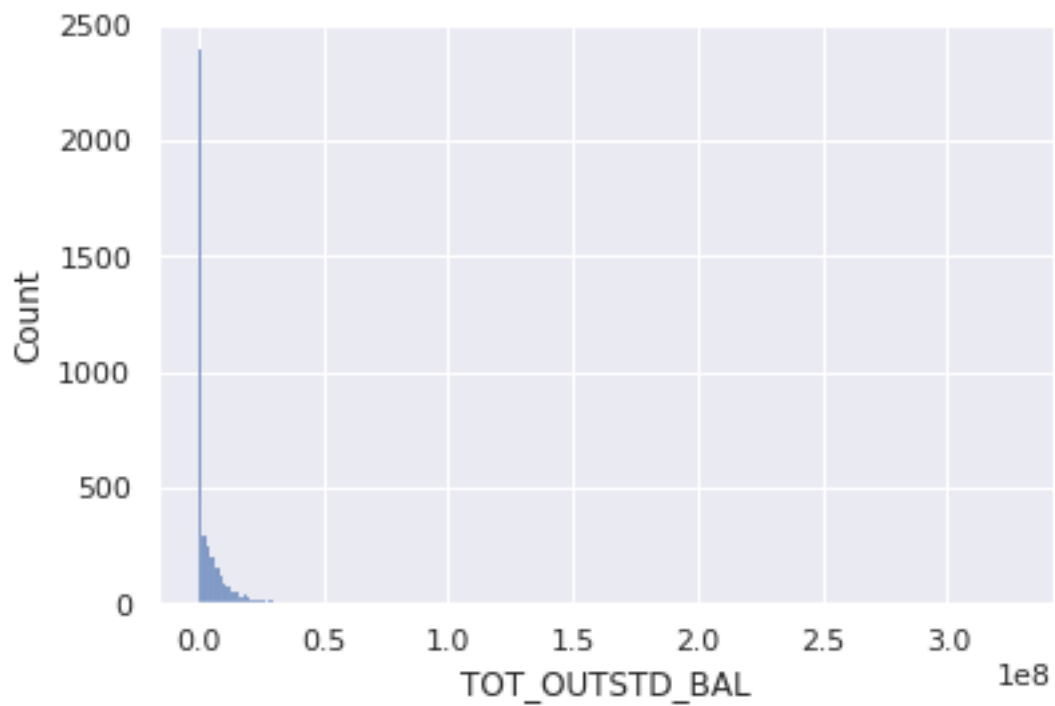
```
This is separate from the ipykernel package so we can avoid doing imports  
until
```

```
[70]: <matplotlib.axes._subplots.AxesSubplot at 0x7fc29d41e550>
```



```
[71]: sns.histplot(df3['TOT_OUTSTD_BAL'])
```

```
[71]: <matplotlib.axes._subplots.AxesSubplot at 0x7fc29df8f3d0>
```



```
[72]: # Applying log transformation to the attributes to normalise the data
```

```
# df3.replace([np.inf, -np.inf], np.nan, inplace=True)
```

```
df3['TOT_OUTSTD_BAL_LOG'] = np.log(df3['TOT_OUTSTD_BAL'])  
sns.histplot(df3['TOT_OUTSTD_BAL_LOG'])
```

```
/usr/local/lib/python3.7/dist-packages/pandas/core/arraylike.py:364:
```

```
RuntimeWarning: divide by zero encountered in log
```

```
    result = getattr(ufunc, method)(*inputs, **kwargs)
```

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:5:
```

```
SettingWithCopyWarning:
```

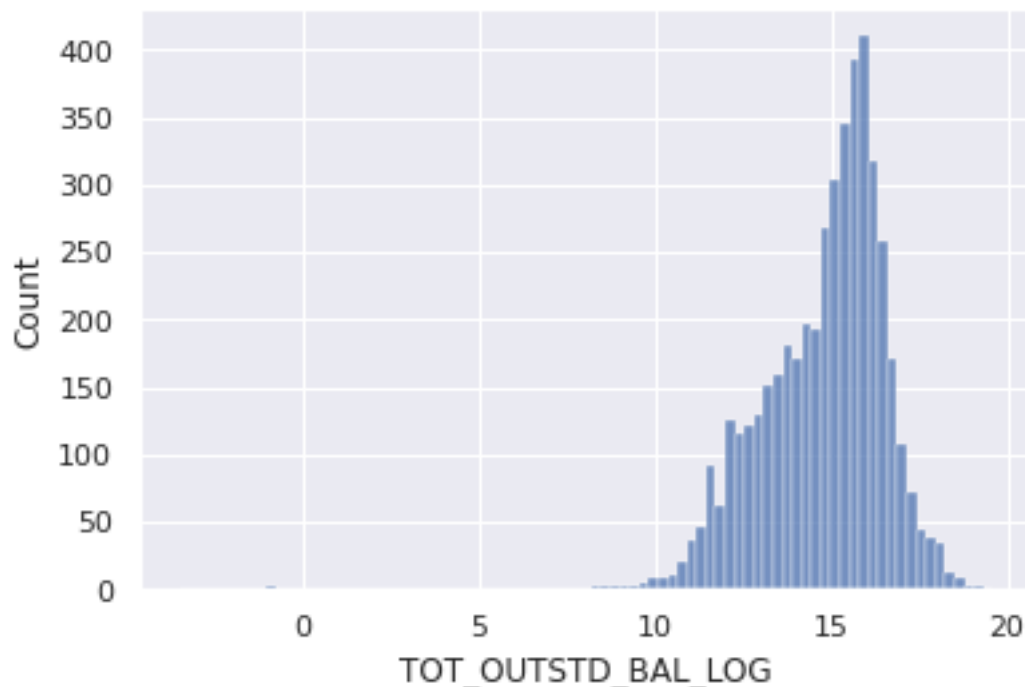
```
A value is trying to be set on a copy of a slice from a DataFrame.
```

```
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
```

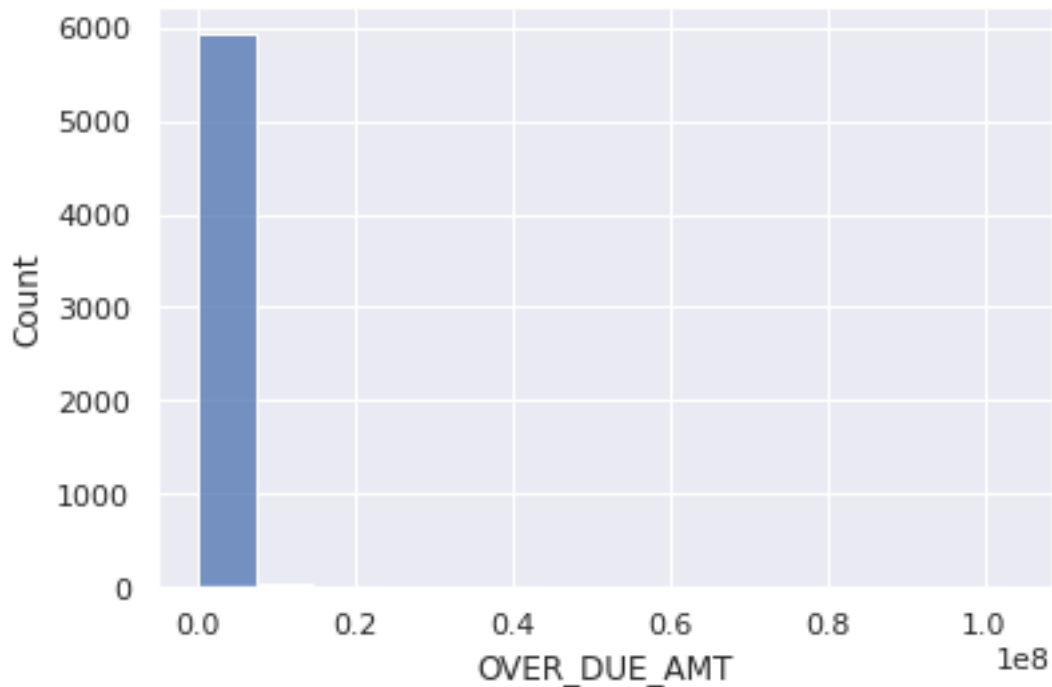
```
"""
```

```
[72]: <matplotlib.axes._subplots.AxesSubplot at 0x7fc29cfe72d0>
```



```
[73]: sns.histplot(df3['OVER_DUE_AMT'])
```

[73]: <matplotlib.axes._subplots.AxesSubplot at 0x7fc29cecf050>



```
[74]: # Applying log transformation to the attributes to normalise the data
```

```
df3['OVER_DUE_AMT_LOG'] = np.log(df3['OVER_DUE_AMT'])
sns.histplot(df3['OVER_DUE_AMT_LOG'])
```

```
/usr/local/lib/python3.7/dist-packages/pandas/core/arraylike.py:364:
```

```
RuntimeWarning: divide by zero encountered in log
```

```
    result = getattr(ufunc, method)(*inputs, **kwargs)
```

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:3:
```

```
SettingWithCopyWarning:
```

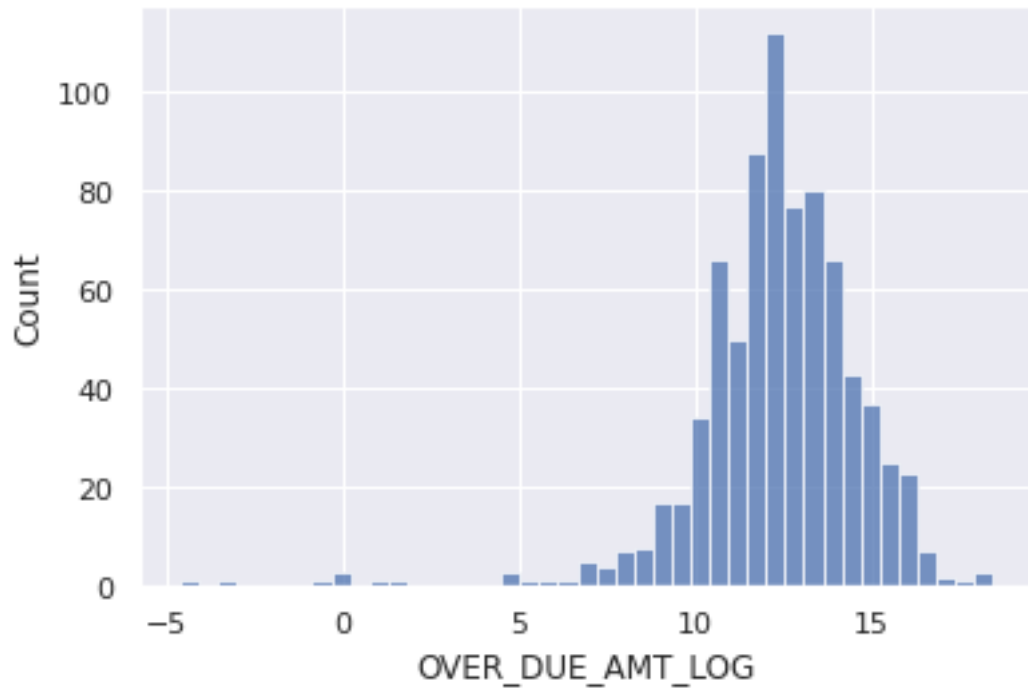
```
A value is trying to be set on a copy of a slice from a DataFrame.
```

```
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

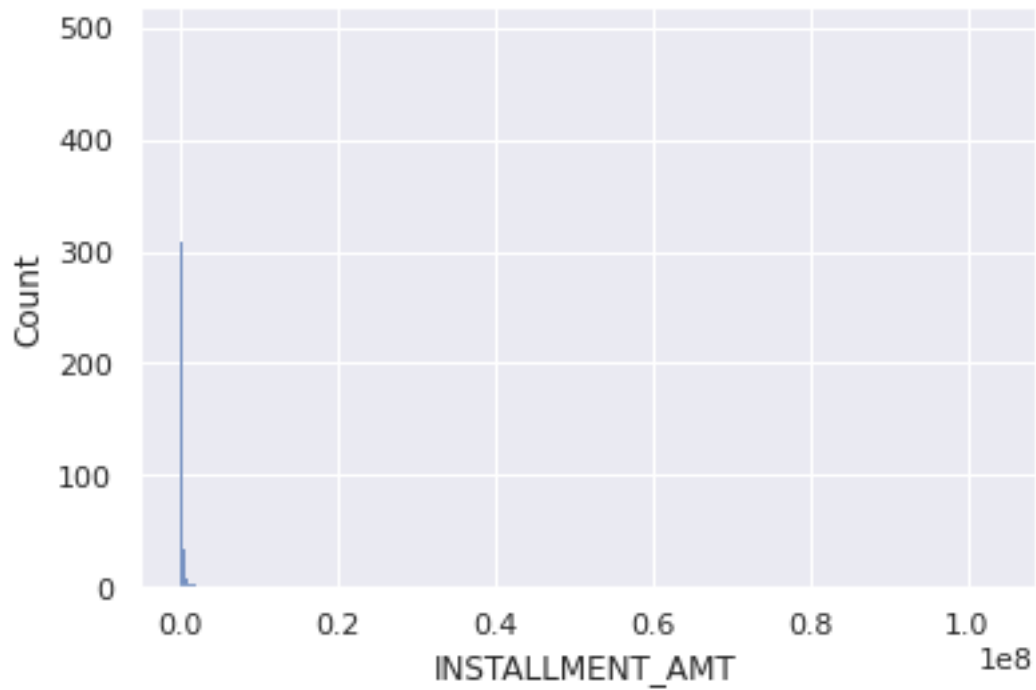
This is separate from the ipykernel package so we can avoid doing imports until

[74]: <matplotlib.axes._subplots.AxesSubplot at 0x7fc29ce643d0>



```
[75]: sns.histplot(df3['INSTALLMENT_AMT'])
```

```
[75]: <matplotlib.axes._subplots.AxesSubplot at 0x7fc29cd52c90>
```



```
[76]: # Applying log transformation to the attributes to normalise the data
```

```
df3['INSTALLMENT_AMT_LOG'] = np.log(df3['INSTALLMENT_AMT'])  
sns.histplot(df3['INSTALLMENT_AMT_LOG'])
```

```
/usr/local/lib/python3.7/dist-packages/pandas/core/arraylike.py:364:
```

```
RuntimeWarning: divide by zero encountered in log
```

```
    result = getattr(ufunc, method)(*inputs, **kwargs)
```

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:3:
```

```
SettingWithCopyWarning:
```

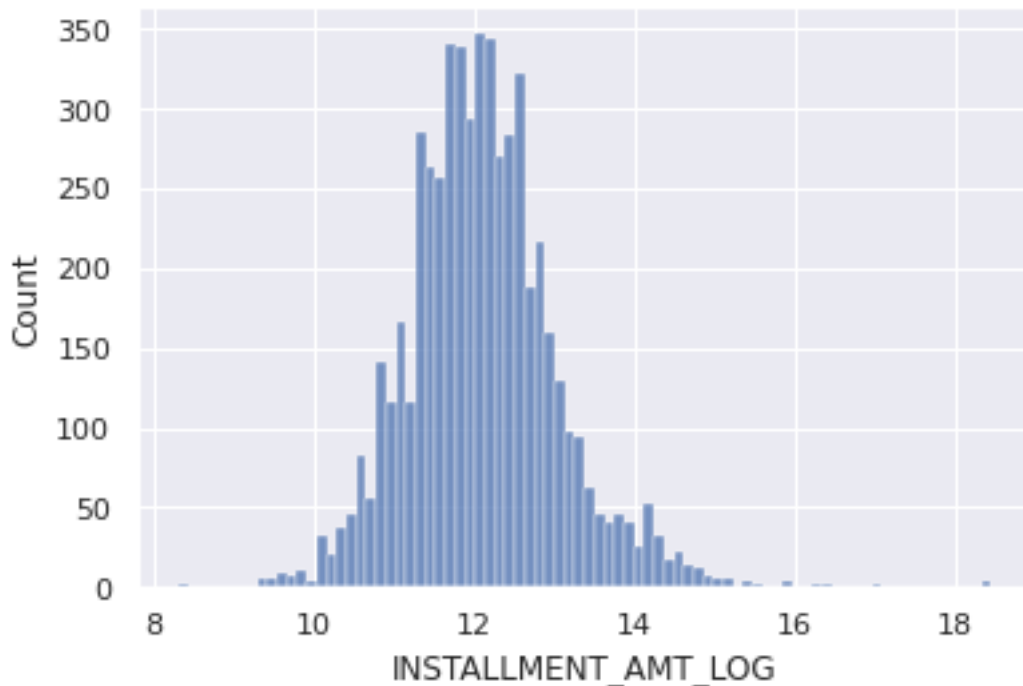
```
A value is trying to be set on a copy of a slice from a DataFrame.
```

```
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

This is separate from the ipykernel package so we can avoid doing imports until

```
[76]: <matplotlib.axes._subplots.AxesSubplot at 0x7fc29cb23e50>
```



```
[77]: # Handle the inifinity values created from log functions.
```

```
df3.replace([np.inf, -np.inf], 0, inplace=True)
df3.head()
```

/usr/local/lib/python3.7/dist-packages/pandas/core/frame.py:5244:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
method=method,

```
[77]:
```

	SUBJECT_ID	ACCOUNT_ID	CURR_CODE	REPAYMENT_FREQUENCY	\
0	110113010000498773	2201160002396407	TZS	MonthlyInstalments30Days	
1	110115010001596750	2201170002990772	TZS	MonthlyInstalments30Days	
2	110114010000832340	2201140001177207	TZS	MonthlyInstalments30Days	
3	110114010000740514	2201160002371670	TZS	MonthlyInstalments30Days	
4	110114010000811717	2201140001161294	TZS	MonthlyInstalments30Days	

	NUMBER_OF_INSTALLMENTS	SANCTION_AMT	TOT_OUTSTD_BAL	OVER_DUE_AMT	\
0	6.0	400000.0	0.00	205609.4	
1	61.0	3500000.0	3272849.49	0.0	
2	57.0	4200000.0	870462.93	0.0	
3	24.0	60000000.0	15740676.93	0.0	
4	69.0	3400000.0	828981.51	0.0	

	INSTALLMENT_AMT	LOAN_TYPE	LOAN_STATUS	LOAN_TENURE	DPD	\
0	0.00	OtherInstalmentOperation	Existing	183.0	399.0	
1	100243.11	ConsumerLoan	Existing	1832.0	0.0	
2	106652.40	ConsumerLoan	Existing	1735.0	0.0	
3	2987342.81	BusinessLoan	Existing	730.0	18.0	
4	99164.00	ConsumerLoan	Existing	2095.0	0.0	

	SANCTION_AMT_LOG	TOT_OUTSTD_BAL_LOG	OVER_DUE_AMT_LOG	INSTALLMENT_AMT_LOG
0	12.899220	0.000000	12.233734	0.000000
1	15.068274	15.001172	0.000000	11.515354
2	15.250595	13.676780	0.000000	11.577330
3	17.909855	16.571759	0.000000	14.909895
4	15.039286	13.627953	0.000000	11.504530

```
[110]: fig = plt.figure(figsize=(16,8))

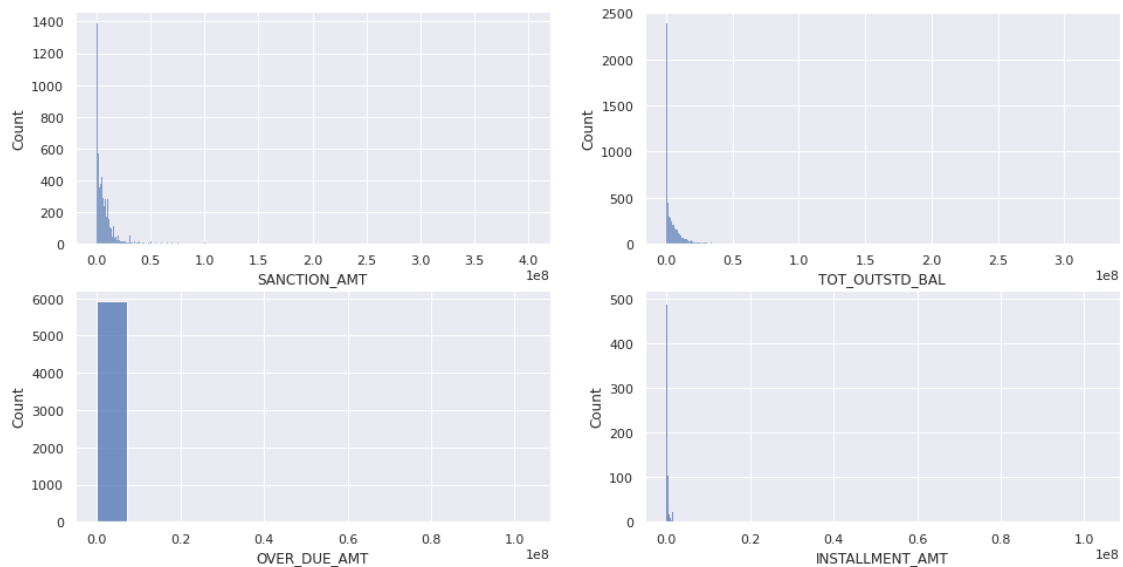
plt.subplot(2, 2, 1)
sns.histplot(df3['SANCTION_AMT'])

plt.subplot(2, 2, 2)
sns.histplot(df3['TOT_OUTSTD_BAL'])
```

```
plt.subplot(2, 2, 3)
sns.histplot(df3['OVER_DUE_AMT'])

plt.subplot(2, 2, 4)
sns.histplot(df3['INSTALLMENT_AMT'])

plt.show()
```



```
[111]: fig = plt.figure(figsize=(16,8))

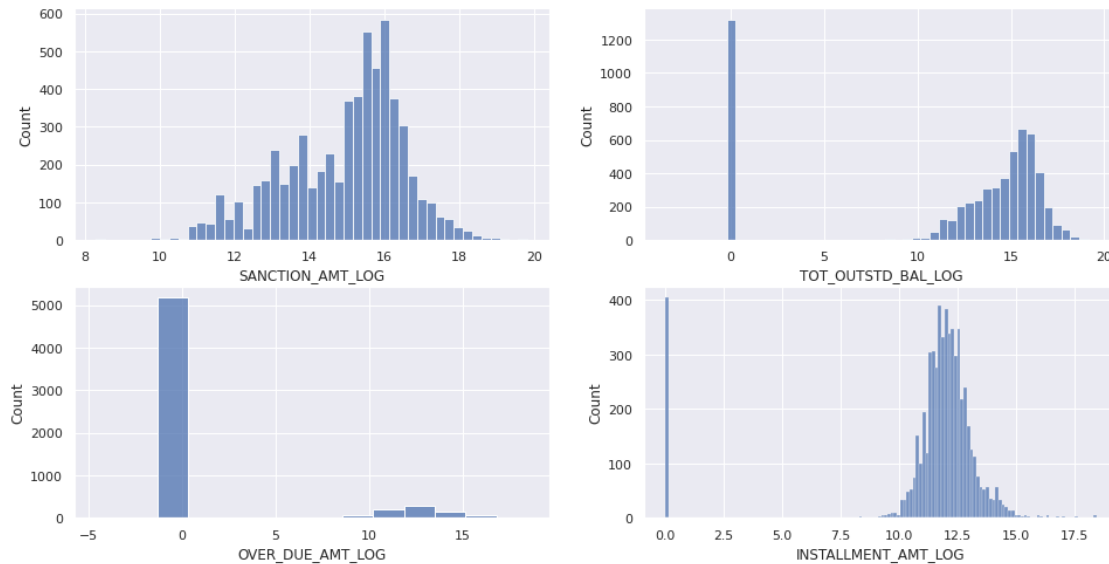
plt.subplot(2, 2, 1)
sns.histplot(df3['SANCTION_AMT_LOG'])

plt.subplot(2, 2, 2)
sns.histplot(df3['TOT_OUTSTD_BAL_LOG'])

plt.subplot(2, 2, 3)
sns.histplot(df3['OVER_DUE_AMT_LOG'])

plt.subplot(2, 2, 4)
sns.histplot(df3['INSTALLMENT_AMT_LOG'])

plt.show()
```

```
[83]: # Create the Risk Identifier Dependant Variable from DPD
      # Logic: A Loan is default if has DPD > 90 days
```

```
df3['IS_DEFAULT'] = np.where(df3['DPD'] >= 90, 1, 0)
df3.head()
```

```
[83]:
```

	SUBJECT_ID	ACCOUNT_ID	CURR_CODE	REPAYMENT_FREQUENCY	\
0	110113010000498773	2201160002396407	TZS	MonthlyInstalments30Days	
1	110115010001596750	2201170002990772	TZS	MonthlyInstalments30Days	
2	110114010000832340	2201140001177207	TZS	MonthlyInstalments30Days	
3	110114010000740514	2201160002371670	TZS	MonthlyInstalments30Days	
4	110114010000811717	2201140001161294	TZS	MonthlyInstalments30Days	

	NUMBER_OF_INSTALLMENTS	SANCTION_AMT	TOT_OUTSTD_BAL	OVER_DUE_AMT	\
0	6.0	400000.0	0.00	205609.4	
1	61.0	3500000.0	3272849.49	0.0	
2	57.0	4200000.0	870462.93	0.0	
3	24.0	60000000.0	15740676.93	0.0	
4	69.0	3400000.0	828981.51	0.0	

	INSTALLMENT_AMT	LOAN_TYPE	LOAN_STATUS	LOAN_TENURE	DPD	\
0	0.00	OtherInstalmentOperation	Existing	183.0	399.0	
1	100243.11	ConsumerLoan	Existing	1832.0	0.0	
2	106652.40	ConsumerLoan	Existing	1735.0	0.0	
3	2987342.81	BusinessLoan	Existing	730.0	18.0	
4	99164.00	ConsumerLoan	Existing	2095.0	0.0	

	SANCTION_AMT_LOG	TOT_OUTSTD_BAL_LOG	OVER_DUE_AMT_LOG	\
--	------------------	--------------------	------------------	---

0	12.899220	0.000000	12.233734
1	15.068274	15.001172	0.000000
2	15.250595	13.676780	0.000000
3	17.909855	16.571759	0.000000
4	15.039286	13.627953	0.000000

	INSTALLMENT_AMT_LOG	IS_DEFAULT
0	0.000000	1
1	11.515354	0
2	11.577330	0
3	14.909895	0
4	11.504530	0

8 Correlation Matrix

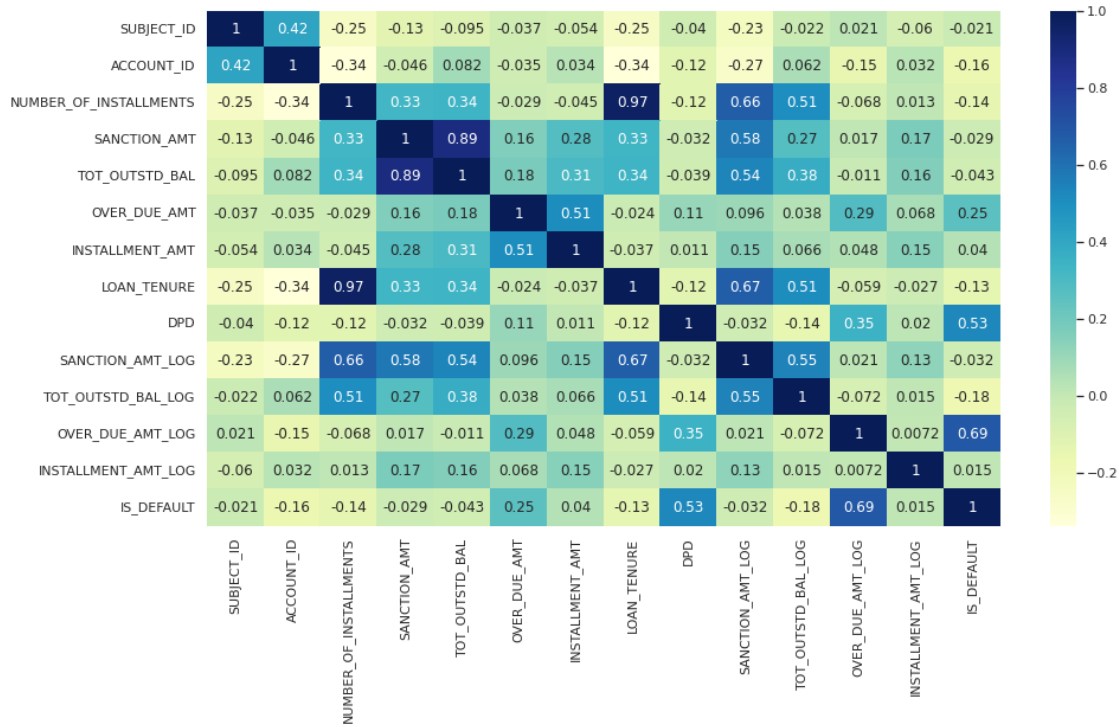
```
[120]: # Keep a copy of the dataframe

df4 = df3.copy()

# Plot the heatmap to see the best correlated variables

corr = df4.corr()
plt.figure(figsize=(15,8))
sns.heatmap(corr, annot=True, cmap='YlGnBu')
```

```
[120]: <matplotlib.axes._subplots.AxesSubplot at 0x7fc289629ed0>
```



```
[121]: # Subset only the relevant variables for the modelling

features = ['REPAYMENT_FREQUENCY', 'NUMBER_OF_INSTALLMENTS', 'SANCTION_AMT_LOG',
            'TOT_OUTSTD_BAL_LOG', 'OVER_DUE_AMT_LOG', 'INSTALLMENT_AMT',
            'LOAN_TYPE', 'LOAN_STATUS', 'LOAN_TENURE', 'IS_DEFAULT']
df4 = df4[features]
df4.head()
```

```
[121]:      REPAYMENT_FREQUENCY  NUMBER_OF_INSTALLMENTS  SANCTION_AMT_LOG \
0  MonthlyInstalments30Days                6.0          12.899220
1  MonthlyInstalments30Days               61.0          15.068274
2  MonthlyInstalments30Days               57.0          15.250595
3  MonthlyInstalments30Days               24.0          17.909855
4  MonthlyInstalments30Days               69.0          15.039286

      TOT_OUTSTD_BAL_LOG  OVER_DUE_AMT_LOG  INSTALLMENT_AMT \
0           0.000000        12.233734           0.00
1          15.001172         0.000000        100243.11
2          13.676780         0.000000        106652.40
3          16.571759         0.000000       2987342.81
4          13.627953         0.000000         99164.00

      LOAN_TYPE  LOAN_STATUS  LOAN_TENURE  IS_DEFAULT
0  OtherInstalmentOperation    Existing        183.0         1
```

1	ConsumerLoan	Existing	1832.0	0
2	ConsumerLoan	Existing	1735.0	0
3	BusinessLoan	Existing	730.0	0
4	ConsumerLoan	Existing	2095.0	0

9 Label Encoding

```
[122]: from sklearn.preprocessing import LabelEncoder

cols = ['REPAYMENT_FREQUENCY', 'LOAN_TYPE', 'LOAN_STATUS']
le = LabelEncoder()
for col in cols:
    df4[col] = le.fit_transform(df4[col])

df4.head()
```

```
[122]:  REPAYMENT_FREQUENCY  NUMBER_OF_INSTALLMENTS  SANCTION_AMT_LOG  \
0                4                6.0          12.899220
1                4               61.0          15.068274
2                4               57.0          15.250595
3                4               24.0          17.909855
4                4               69.0          15.039286

    TOT_OUTSTD_BAL_LOG  OVER_DUE_AMT_LOG  INSTALLMENT_AMT  LOAN_TYPE  \
0          0.000000      12.233734          0.00          5
1          15.001172       0.000000        100243.11          1
2          13.676780       0.000000        106652.40          1
3          16.571759       0.000000        2987342.81          0
4          13.627953       0.000000         99164.00          1

    LOAN_STATUS  LOAN_TENURE  IS_DEFAULT
0            0        183.0           1
1            0       1832.0           0
2            0       1735.0           0
3            0        730.0           0
4            0       2095.0           0
```

10 Train Test Split

```
[123]: # Specifying the input and output attributes

X = df4.drop(columns=['IS_DEFAULT'], axis=1)
y = df4['IS_DEFAULT']
```

```
[124]: from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25,
↳random_state=42)
```

11 Model Training

```
[125]: # Classify function
from sklearn.model_selection import cross_val_score

def classify_model(model, x, y):
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25,
↳random_state=42)
    model.fit(X_train, y_train)
    print('Accuracy is: ', model.score(X_test, y_test)*100)

    # Cross validation
    score = cross_val_score(model, X, y, cv=5)
    print('Cross Validation Score: ', np.mean(score)*100)
```

```
[126]: # Logistic Regression Model
from sklearn.linear_model import LogisticRegression

model = LogisticRegression()
classify_model(model, X, y)
```

Accuracy is: 93.8255033557047
Cross Validation Score: 94.12751677852349

```
[127]: # DecisionTree Classifier Model

from sklearn.tree import DecisionTreeClassifier

model = DecisionTreeClassifier()
classify_model(model, X, y)
```

Accuracy is: 97.71812080536913
Cross Validation Score: 97.09731543624163

```
[128]: # RandomForest Classifier Model

from sklearn.ensemble import RandomForestClassifier

model = RandomForestClassifier()
classify_model(model, X, y)
```

Accuracy is: 98.45637583892616
Cross Validation Score: 98.02013422818791

12 Confusion Matrix

```
[129]: # Since the RandomForest Model perfoms better, we shall use it further
```

```
model = RandomForestClassifier()  
model.fit(X_train, y_train)
```

```
[129]: RandomForestClassifier()
```

```
[130]: from sklearn.metrics import confusion_matrix
```

```
y_pred = model.predict(X_test)  
cm = confusion_matrix(y_test, y_pred)  
cm
```

```
[130]: array([[1379,  19],  
          [  8,  84]])
```

```
[131]: sns.heatmap(cm, annot=True, cmap='BrBG_r')
```

```
[131]: <matplotlib.axes._subplots.AxesSubplot at 0x7fc289406cd0>
```

